


Adversarial classification: can a classifier tell synthetic from real?

Here the diagonal is the GOOD result -- it means the two are indistinguishable



Note the inverted reading. In prediction tasks $AUC \rightarrow 1.0$ is success; in synthetic-data validation $AUC \rightarrow 0.50$ is success. A curve bowing toward the top-left would mean the synthetic records carry detectable artefacts.